

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from datetime import datetime
pd.options.display.float_format = '{:.2f}'.format

from itertools import combinations
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.stattools import acf, pacf
from statsmodels.tsa.arima_model import ARIMA as ARIMA
import statsmodels.api as sm
import statsmodels.tsa.api as smt
```

In [2]:

```
data = pd.read_csv('AirPassengers.csv')
```

In [3]:

```
data.head()
```

Out[3]:

	Month	#Passengers
0	1949-01	112
1	1949-02	118
2	1949-03	132
3	1949-04	129
4	1949-05	121

In [4]:

```
data.shape
```

Out[4]:

```
(144, 2)
```

In [5]:

```
data.columns
```

Out[5]:

```
Index(['Month', '#Passengers'], dtype='object')
```

In [6]:

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144 entries, 0 to 143
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Month            144 non-null    object
1   #Passengers      144 non-null    int64
dtypes: int64(1), object(1)
memory usage: 2.4+ KB
```

In [10]:

data.describe()

Out[10]:

#Passengers	
count	144.00
mean	280.30
std	119.97
min	104.00
25%	180.00
50%	265.50
75%	360.50
max	622.00

In [11]:

```
data['Date'] = pd.to_datetime(data['Month'])
data = data.drop(columns = 'Month')
data = data.set_index('Date')
data = data.rename(columns = {'#Passengers': 'Passengers'})
data.head()
```

Out[11]:

Passengers	
Date	
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121


In [12]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 144 entries, 1949-01-01 to 1960-12-01
Data columns (total 1 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Passengers  144 non-null    int64
dtypes: int64(1)
memory usage: 2.2 KB
```


In [13]:

```
plt.figure(figsize = (15,5))
data['Passengers'].plot();
```



In [14]:


```
dec = sm.tsa.seasonal_decompose(data['Passengers'],period = 12, model = 'multiplicative').p
plt.show()
```



In [16]:

```
data_diff = data.diff()
data_diff = data_diff.dropna()

dec = sm.tsa.seasonal_decompose(data_diff, period = 12).plot()
plt.show()
```



In [17]:

```
def test_stationarity(timeseries):
    #Determining rolling statistics
    MA = timeseries.rolling(window=12).mean()
    MSTD = timeseries.rolling(window=12).std()

    #Plot rolling statistics:
    plt.figure(figsize=(15,5))
    orig = plt.plot(timeseries, color='blue',label='Original')
    mean = plt.plot(MA, color='red', label='Rolling Mean')
    std = plt.plot(MSTD, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show(block=False)

    #Perform Dickey-Fuller test:
    print('Results of Dickey-Fuller Test:')
    dfctest = adfuller(timeseries, autolag='AIC')
    dfcoutput = pd.Series(dfctest[0:4], index=['Test Statistic','p-value','#Lags Used','Number of Observations Used'])
    for key,value in dfctest[4].items():
        dfcoutput['Critical Value (%s)'%key] = value
    print(dfcoutput)
```

In [18]:


```
def tsplot(y, lags=None, figsize=(12, 7), style='bmh'):
    if not isinstance(y, pd.Series):
        y = pd.Series(y)

    with plt.style.context(style):
        fig = plt.figure(figsize=figsize)
        layout = (2, 2)
        ts_ax = plt.subplot2grid(layout, (0, 0), colspan=2)
        acf_ax = plt.subplot2grid(layout, (1, 0))
        pacf_ax = plt.subplot2grid(layout, (1, 1))

        y.plot(ax=ts_ax)
        p_value = sm.tsa.stattools.adfuller(y)[1]
        ts_ax.set_title('Time Series Analysis Plots\n Dickey-Fuller: p={0:.5f}'.format(p_value))
        smt.graphics.plot_acf(y, lags=lags, ax=acf_ax)
        smt.graphics.plot_pacf(y, lags=lags, ax=pacf_ax)
        plt.tight_layout()
```


In [19]:

```
plt.figure(figsize = (15,5))
data['Passengers'].plot();
```




In [20]:

```
dec = sm.tsa.seasonal_decompose(data['Passengers'], period = 12, model = 'multiplicative').p
plt.show()
```



In [21]:

```
test_stationarity(data['Passengers'])
```




Results of Dickey-Fuller Test:

Test Statistic	0.82
p-value	0.99
#Lags Used	13.00
Number of Observations Used	130.00
Critical Value (1%)	-3.48
Critical Value (5%)	-2.88
Critical Value (10%)	-2.58

dtype: float64


In [22]:

```
data_diff = data.diff()  
data_diff = data_diff.dropna()  
  
dec = sm.tsa.seasonal_decompose(data_diff, period = 12).plot()  
plt.show()
```



In [23]:

```
test_stationarity(data_diff)
```




Results of Dickey-Fuller Test:

Test Statistic	-2.83
p-value	0.05
#Lags Used	12.00
Number of Observations Used	130.00
Critical Value (1%)	-3.48
Critical Value (5%)	-2.88
Critical Value (10%)	-2.58
dtype:	float64


In [24]:

```
tsplot(data[ 'Passengers' ])
```



In [25]:

```
tsplot(data_diff['Passengers'])
```



In [26]:

```
model = ARIMA(data['Passengers'], order = (2,1,2))
model_fit = model.fit()
print(model_fit.summary())
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py:47

2: FutureWarning:

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)
```

```
warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
warnings.warn('No frequency information was'
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
warnings.warn('No frequency information was'
```

ARIMA Model Results

```
=====
=====
Dep. Variable:          D.Passengers   No. Observations:
143
Model:                  ARIMA(2, 1, 2)   Log Likelihood          -66
6.022
Method:                  css-mle         S.D. of innovations         2
4.714
Date:                   Wed, 14 Sep 2022   AIC                        134
4.043
Time:                   10:06:41          BIC                        136
1.820
Sample:                 02-01-1949        HQIC                       135
1.267
                        - 12-01-1960
=====
=====
                        coef      std err          z      P>|z|      [0.025
0.975]
-----
const                2.5311        0.708        3.574      0.000        1.143
3.919
ar.L1.D.Passengers    1.6477        0.033       49.933      0.000        1.583
1.712
```

ar.L2.D.Passengers	-0.9094	0.033	-27.880	0.000	-0.973
-0.845					
ma.L1.D.Passengers	-1.9098	0.065	-29.515	0.000	-2.037
-1.783					
ma.L2.D.Passengers	0.9997	0.068	14.809	0.000	0.867
1.132					

Roots

=====				
===				
	Real	Imaginary	Modulus	Freque
ncy				

AR.1	0.9059	-0.5281j	1.0486	-0.0
840				
AR.2	0.9059	+0.5281j	1.0486	0.0
840				
MA.1	0.9552	-0.2965j	1.0002	-0.0
479				
MA.2	0.9552	+0.2965j	1.0002	0.0
479				

In [27]:

```

size = int(len(data) - 30)
train, test = data['Passengers'][0:size], data['Passengers'][size:len(data)]

print('\t ARIMA MODEL : In- Sample Forecasting \n')

history = [x for x in train]
predictions = []

for t in range(len(test)):

    model = ARIMA(history, order=(2,1,2))
    model_fit = model.fit(dispatch = 0)

    output = model_fit.forecast()
    yhat = output[0]
    predictions.append(float(yhat))

    obs = test[t]
    history.append(obs)

    print('predicted = %f, expected = %f' % (yhat, obs))

```

ARIMA MODEL : In- Sample Forecasting

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py:47

2: FutureWarning:

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```

import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)

warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)

```

```

predicted = 433.256972, expected = 491.000000
predicted = 478.355854, expected = 505.000000
predicted = 474.552701, expected = 404.000000
predicted = 367.687051, expected = 359.000000
predicted = 386.046122, expected = 310.000000
predicted = 300.551728, expected = 337.000000
predicted = 342.709246, expected = 360.000000
predicted = 374.434495, expected = 342.000000
predicted = 368.418676, expected = 406.000000
predicted = 427.293772, expected = 396.000000
predicted = 416.580357, expected = 420.000000
predicted = 431.952427, expected = 472.000000
predicted = 465.574801, expected = 548.000000

```

```

predicted = 516.133868, expected = 559.000000
predicted = 522.642349, expected = 463.000000
predicted = 407.122831, expected = 407.000000
predicted = 367.581910, expected = 362.000000
predicted = 349.941636, expected = 405.000000
predicted = 415.817585, expected = 417.000000
predicted = 443.407937, expected = 391.000000
predicted = 432.877393, expected = 419.000000
predicted = 467.788485, expected = 461.000000
predicted = 505.289402, expected = 472.000000
predicted = 505.208366, expected = 535.000000
predicted = 548.678029, expected = 622.000000
predicted = 603.217422, expected = 606.000000
predicted = 560.803257, expected = 508.000000
predicted = 458.420863, expected = 461.000000
predicted = 419.481597, expected = 390.000000
predicted = 373.834768, expected = 432.000000

```


In [28]:

```

predictions_series = pd.Series(predictions, index = test.index)
fig,ax = plt.subplots(nrows = 1,ncols = 1,figsize = (15,5))

plt.subplot(1,1,1)
plt.plot(data['Passengers'],label = 'Expected Values')
plt.plot(predictions_series,label = 'Predicted Values');
plt.legend(loc="upper left")
plt.show()

```



In [29]:

```

error = np.sqrt(mean_squared_error(test,predictions))
print('Test RMSE: %.4f' % error)

```

Test RMSE: 42.5173

In [31]:

```

#our program ends here
#another method
#out of sampling forecasting

```

In [32]:

```
from pandas.tseries.offsets import DateOffset
future_dates = [data.index[-1] + DateOffset(weeks = x) for x in range(0,49)]

# New dataframe for storing the future values
df1 = pd.DataFrame(index = future_dates[1:],columns = data.columns)

forecast = pd.concat([data,df1])
forecast['ARIMA_Forecast_Function'] = np.NaN
forecast['ARIMA_Predict_Function'] = np.NaN
forecast.head()
```

Out[32]:

	Passengers	ARIMA_Forecast_Function	ARIMA_Predict_Function
1949-01-01	112	NaN	NaN
1949-02-01	118	NaN	NaN
1949-03-01	132	NaN	NaN
1949-04-01	129	NaN	NaN
1949-05-01	121	NaN	NaN

In [33]:

```

ARIMA_history_f = [x for x in train]
f1 = []

for t in range(len(df1)):

    model = ARIMA(ARIMA_history_f, order = (2,1,2))
    model_fit = model.fit(dispatch=0)

    output = model_fit.forecast()[0][0]

    ARIMA_history_f.append(output)
    f1.append(output)

for i in range(len(f1)):
    forecast.iloc[144 + i,1] = f1[i]
forecast.tail()

```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py:47

2: FutureWarning:

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```

import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)

warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)


```

Out[33]:

	Passengers	ARIMA_Forecast_Function	ARIMA_Predict_Function
1961-10-05	NaN	490.80	NaN
1961-10-12	NaN	493.31	NaN
1961-10-19	NaN	495.81	NaN
1961-10-26	NaN	498.32	NaN
1961-11-02	NaN	500.83	NaN

In [34]:

```
forecast[['Passengers', 'ARIMA_Forecast_Function']].plot(figsize = (12,8));
```



In [35]:

```

ARIMA_history_p = [x for x in train]
f2 = []

for t in range(len(df1)):

    model = ARIMA(ARIMA_history_p, order = (2,1,2))
    model_fit = model.fit(dispatch=0)

    output = model_fit.predict(start = len(ARIMA_history_p),end = len(ARIMA_history_p),typ

    ARIMA_history_p.append(output)
    f2.append(output)

for i in range(len(f2)):
    forecast.iloc[144 + i,2] = f2[i]
forecast.tail()

```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py:47

2: FutureWarning:

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```

import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)

warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)

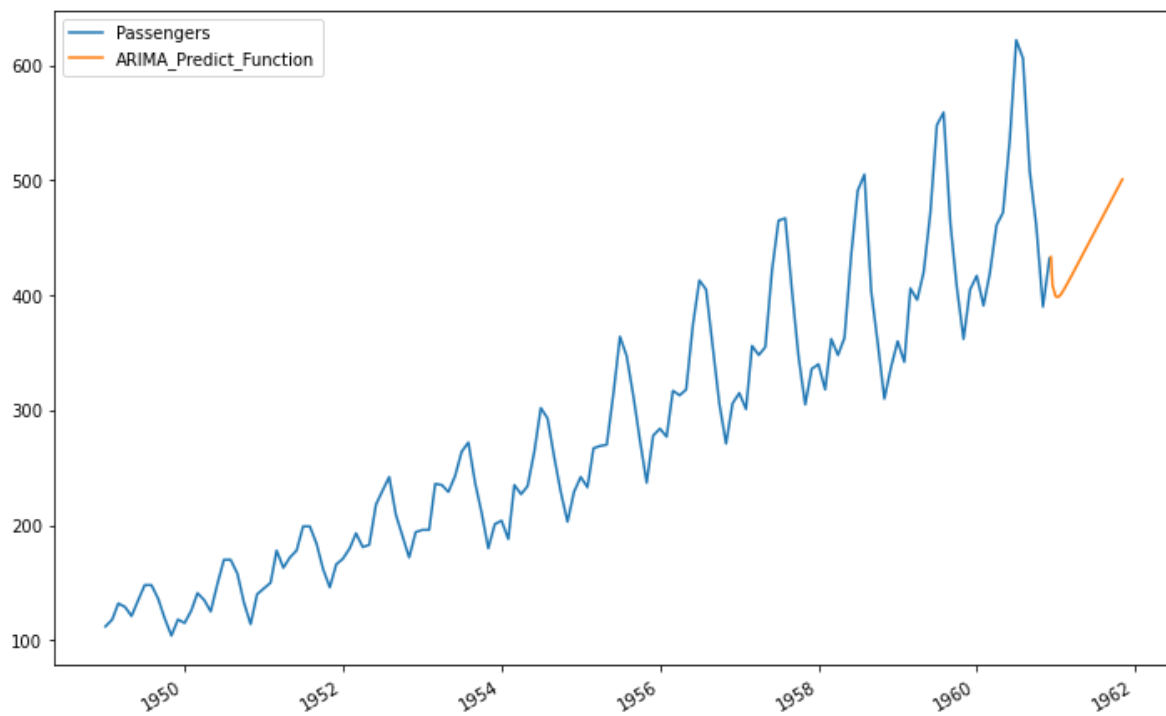
```

Out[35]:

	Passengers	ARIMA_Forecast_Function	ARIMA_Predict_Function
1961-10-05	NaN	490.80	490.80
1961-10-12	NaN	493.31	493.31
1961-10-19	NaN	495.81	495.81
1961-10-26	NaN	498.32	498.32
1961-11-02	NaN	500.83	500.83

In [36]:

```
forecast[['Passengers', 'ARIMA_Predict_Function']].plot(figsize = (12,8));
```



In [37]:

```
sum(f1) == sum(f2)
```

Out[37]:


True

In [38]:

```
#sarimax model  
#arima model ends here
```


In [39]:

```
data_diff_seas = data_diff.diff(12)
data_diff_seas = data_diff_seas.dropna()
dec = sm.tsa.seasonal_decompose(data_diff_seas, period = 12)
dec.plot()
plt.show()
```



In [40]:

```
test_stationarity(data_diff_seas['Passengers'])
```




Results of Dickey-Fuller Test:

Test Statistic	-15.60
p-value	0.00
#Lags Used	0.00
Number of Observations Used	130.00
Critical Value (1%)	-3.48
Critical Value (5%)	-2.88
Critical Value (10%)	-2.58

dtype: float64

In [41]:

```
tsplot(data_diff_seas['Passengers'])
```



In [42]:

```

model = sm.tsa.statespace.SARIMAX(data['Passengers'], order = (2,1,2), seasonal_order = (0,1,
model_fit = model.fit()
print(model_fit.summary())

```

```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.p
y:524: ValueWarning: No frequency information was provided, so inferred freq
uency MS will be used.
  warnings.warn('No frequency information was'
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.p
y:524: ValueWarning: No frequency information was provided, so inferred freq
uency MS will be used.
  warnings.warn('No frequency information was'

```

SARIMAX Results

```

=====
=====
Dep. Variable:          Passengers    No. Observations:
144
Model:                SARIMAX(2, 1, 2)x(0, 1, [1], 12)    Log Likelihood
-503.968
Date:                  Wed, 14 Sep 2022    AIC
1019.935
Time:                  10:14:46    BIC
1037.186
Sample:                01-01-1949    HQIC
1026.945
                        - 12-01-1960
Covariance Type:                opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						

ar.L1	0.3966	0.422	0.940	0.347	-0.430	
1.223						
ar.L2	0.3538	0.317	1.115	0.265	-0.268	
0.976						
ma.L1	-0.7648	0.432	-1.769	0.077	-1.612	
0.083						
ma.L2	-0.2060	0.414	-0.497	0.619	-1.018	
0.606						
ma.S.L12	-0.1033	0.112	-0.921	0.357	-0.323	
0.117						
sigma2	127.2934	14.232	8.944	0.000	99.400	15
5.187						
=====						
=====						
Ljung-Box (L1) (Q):	0.02		Jarque-Bera (JB):			
11.10						
Prob(Q):	0.88		Prob(JB):			
0.00						
Heteroskedasticity (H):	2.57		Skew:			
0.05						
Prob(H) (two-sided):	0.00		Kurtosis:			
4.42						

```
=====
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

In [43]:

```

size = int(len(data) - 30)
train, test = data['Passengers'][0:size], data['Passengers'][size:len(data)]

print('\t SARIMA MODEL : In - Sample Forecasting \n')

history = [x for x in train]
predictions = []

for t in range(len(test)):

    model = sm.tsa.statespace.SARIMAX(history, order = (2,1,2), seasonal_order = (0,1,1,12))
    model_fit = model.fit()

    output = model_fit.forecast()
    yhat = output[0]
    predictions.append(float(yhat))

    obs = test[t]
    history.append(obs)

    print('predicted = %f, expected = %f' % (yhat, obs))

```

SARIMA MODEL : In - Sample Forecasting

```

predicted = 479.084514, expected = 491.000000
predicted = 490.553512, expected = 505.000000
predicted = 441.276126, expected = 404.000000
predicted = 357.270516, expected = 359.000000
predicted = 315.251199, expected = 310.000000
predicted = 347.831240, expected = 337.000000

```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

warnings.warn("Maximum Likelihood optimization failed to "

```


predicted = 344.251813, expected = 360.000000
predicted = 336.835601, expected = 342.000000
predicted = 387.593322, expected = 406.000000
predicted = 387.333485, expected = 396.000000
predicted = 408.192790, expected = 420.000000
predicted = 485.988157, expected = 472.000000
predicted = 529.031342, expected = 548.000000
predicted = 551.914007, expected = 559.000000
predicted = 459.061267, expected = 463.000000
predicted = 411.970100, expected = 407.000000
predicted = 358.421156, expected = 362.000000
predicted = 384.945724, expected = 405.000000
predicted = 420.143813, expected = 417.000000
predicted = 397.755392, expected = 391.000000
predicted = 451.335510, expected = 419.000000
predicted = 415.675504, expected = 461.000000
predicted = 465.295980, expected = 472.000000
predicted = 529.835404, expected = 535.000000
predicted = 599.299659, expected = 622.000000
predicted = 626.292199, expected = 606.000000
predicted = 513.891979, expected = 508.000000
predicted = 450.136741, expected = 461.000000

```

```
predicted = 411.653929, expected = 390.000000  
predicted = 438.411433, expected = 432.000000
```

In [44]:

```
predictions_series = pd.Series(predictions, index = test.index)  
fig,ax = plt.subplots(nrows = 1,ncols = 1,figsize = (15,5))  
  
plt.subplot(1,1,1)  
plt.plot(data['Passengers'],label = 'Expected Values')  
plt.plot(predictions_series,label = 'Predicted Values');  
plt.legend(loc="upper left")  
plt.show()
```



In [45]:

```
error = np.sqrt(mean_squared_error(test,predictions))  
print('Test RMSE: %.4f' % error)
```

Test RMSE: 16.9251

In [46]:

```
#so as we can see , using arima the rmse was 42.5173 but using sarimax you can see the  
#the rmse is 16.9251 , just awesome
```

In []: